NovaMind AI — Products & Pricing

Product Suite Overview

NovaMind AI offers three core product lines: the **NovaMind Core API** (developer-facing model access), **NovaMind Studio** (no-code/low-code AI builder), and **NovaMind AgentOS** (autonomous multi-agent orchestration platform). All products are available via cloud SaaS or private cloud deployment.

1. NovaMind Core API

Programmatic access to NovaMind's foundation models via REST API. Suitable for engineering teams building custom AI applications, chatbots, document intelligence tools, and more.

Available Models:

- **Nova-7B** – Lightweight 7B parameter model. Optimized for low-latency tasks (classification, summarization, Q&A). Context window: 32K tokens.
- **Nova-70B** – High-capability 70B parameter model. Best for complex reasoning, code generation, and long-form content. Context window: 128K tokens.
- **Nova-Vision** – Multimodal model supporting text + image inputs. Ideal for document parsing, chart analysis, and visual Q&A. Context window: 64K tokens.
- **Nova-Embed** – Text embedding model optimized for semantic search and RAG pipelines. 1536-dimensional vectors.

Pricing (pay-as-you-go):

- Nova-7B: $0.12 / 1M input tokens | $0.36 / 1M output tokens
- Nova-70B: $0.80 / 1M input tokens | $2.40 / 1M output tokens
- Nova-Vision: $1.20 / 1M input tokens | $3.60 / 1M output tokens
- Nova-Embed: $0.04 / 1M tokens

Committed Use Discounts: 15% off for 12-month commitments; 25% off for 24-month commitments. Volume discounts available above 500M tokens/month.

2. NovaMind Studio

A visual, no-code/low-code platform for building AI-powered workflows, chatbots, document processors, and dashboards. Designed for business analysts, operations teams, and product teams without deep ML expertise.

Key Features:

- Drag-and-drop workflow builder with 80+ pre-built AI action blocks
- Native connectors: Salesforce, HubSpot, Slack, Notion, SharePoint, Google Drive, Snowflake
- RAG pipeline builder with automatic document chunking and vector indexing
- Built-in evaluation dashboard (accuracy, hallucination rate, latency metrics)
- Role-based access control and audit logging
- SSO via SAML 2.0 / OIDC

Pricing Tiers:

- **Starter – $299/month:** Up to 5 users, 50 workflows, 10K API calls/month. Community support.
- **Growth – $999/month:** Up to 25 users, unlimited workflows, 100K API calls/month. Email + chat support, custom branding.
- **Business – $3,499/month:** Up to 100 users, SSO, dedicated success manager, SLA 99.9%, 500K API calls/month.
- **Enterprise – Custom pricing:** Unlimited users, on-premise or VPC deployment, custom model fine-tuning, 24/7 support, SLA 99.99%.

All plans include a 14-day free trial. Annual billing provides a 20% discount over monthly rates.

3. NovaMind AgentOS

An enterprise-grade platform for deploying, managing, and monitoring autonomous AI agents. AgentOS enables multi-step agentic workflows that can browse the web, execute code, query databases, call APIs, and coordinate between specialized sub-agents.

Core Capabilities:

- Multi-agent orchestration with configurable agent roles and tool access
- Human-in-the-loop approval gates for high-risk actions
- Persistent agent memory (short-term, long-term, episodic)
- Full observability: trace viewer, replay debugging, cost attribution per agent
- Compliance mode: all agent actions logged and exportable for audit
- Pre-built agent templates for: financial research, legal contract review, IT helpdesk automation, sales intelligence

Pricing:

- **Team – $1,999/month:** Up to 10 concurrent agents, 20 users, 200K agent actions/month.
- **Professional – $5,999/month:** Up to 50 concurrent agents, 75 users, 1M agent actions/month, dedicated infra.
- **Enterprise – Custom:** Unlimited agents, on-prem/VPC, custom SLA, professional services included.

Professional Services & Add-Ons

- **Model Fine-Tuning:** Starting at $15,000 per fine-tuning project (includes data prep, training, and evaluation report).
- **RAG Implementation Package:** $8,500 flat fee – NovaMind engineers design and deploy a production RAG pipeline in your environment.
- **AI Readiness Assessment:** $4,000 – 2-week engagement to evaluate your data, infrastructure, and use-case fit.
- **Ongoing ML Ops Support:** $3,500/month retainer – dedicated ML engineer support for model monitoring and iteration.

SLAs & Support

- API uptime SLA: 99.9% (Growth/Business), 99.99% (Enterprise)

- Support response times: Community (48h), Email/Chat (4h), Business (1h), Enterprise (15 min)
- Status page: status.novamind.ai